

AI Agents for Business Value and Responsible Governance

Audience: Business development leaders, governance teams, and non-technical decision-makers | **Format:** Explain → demonstrate → practise → govern → decide | **Facilitator:** Nicodemus Chan

INDEPENDENT TEACHING SAMPLE: This lesson plan is an independent demonstration by Nicodemus Chan; it is not official OpenAI curriculum.

OPENING INCIDENT (0–5 min): The business wants an agent that answers policy questions and can approve expenses. Its search tool finds the relevant travel policy, but the requested daily meal limit is absent. Should the agent answer, act, or escalate—and who owns the risk?

Five measurable learning outcomes

1. Distinguish an agent from a chatbot or fixed automation by identifying who controls the workflow and chooses the next step.
2. Use three business-fit criteria—complex decisions, brittle rules, and unstructured data—to accept or reject a proposed agent use case.
3. Map a candidate agent to its model, tools, and instructions, and classify tools as data, action, or orchestration.
4. Select a single-agent or multi-agent pattern and justify the choice using workflow complexity, tool overlap, evaluation, and maintenance.
5. Define a minimum governance plan with layered guardrails, tool-risk controls, exit conditions, and human-intervention triggers.

Session sequence and formative checks

TIME	STAGE	FACILITATOR + PARTICIPANT ACTIVITY	FORMATIVE KNOWLEDGE CHECK
0–5	Opening / predict	Choose: answer, act, or escalate before the policy-search result is explained.	What business harm could follow from each choice?
5–14	Explain: agent fit	Use slides 2–4: agent vs chatbot; suitable workflows have complex decisions, brittle rules, or unstructured data.	Who chooses the next step—the user, fixed code, or the model? Does this case need an agent?
14–23	Explain: foundations	Use slides 5–10: model tradeoffs; model + tools + instructions; data, action, and orchestration tools; clear edge cases.	Classify CRM read, CRM update, and specialist delegation. Which tool creates the greatest risk?
23–32	Demonstrate	Show the policy agent trace: question → search tool → relevant policy → requested amount absent → answer withheld and escalation offered.	Did retrieval authorize an expense action? Which instruction or guardrail prevents it?
32–44	Participant practice	Teams map one business workflow: goal, agent-fit criterion, data/action tools, instructions, owner, risks, and success evidence.	Each team must name one action requiring approval and one measurable pilot outcome.
44–51	Retrieval diagnostic	Diagnose the policy trace as tool failure, missing data, instruction failure, or governance failure; choose the safe next step.	What evidence is missing? Should the agent retry, abstain, ask, or hand control to a person?
51–56	Orchestration + guardrails	Start with one agent; add specialists only for complex logic or persistent tool-selection failure. Layer controls and human intervention.	Why not begin with multiple agents? When must a person approve or take over?
56–60	Assessment / exit	Individual response to the governance decision below; sample two recommendations aloud.	Success requires use-case fit, minimum design, control owner, and an evidence-based pilot decision.

Retrieval-failure diagnostic exercise

BUSINESS DECISION TASK: Required trace: the data tool finds travel_policy.txt, but no daily meal limit. The agent also has an expense-approval action tool. Teams identify the missing evidence, block the action, choose abstain/clarify/escalate, name the accountable human, and add the case to evaluation. Expected diagnosis: retrieval is relevant but insufficient; no financial action is authorized.

MATERIALS + PREREQUISITES: Nicodemus_Chan_AI_Agents_Teaching_Demo.pptx; OpenAI’s A practical guide to building agents; agent-use-case canvas; scenario and risk cards; policy-search demonstration or saved transcript; voting cards. No coding is required. Participants should know one candidate workflow, its decision owner, and the organization’s risk tolerance.

COMMON PROBLEMS + FALLBACKS: Video or internet unavailable → use still images and the slide narrative. Live agent, credentials, or quota unavailable → use the saved policy trace. Tool integrations unavailable → role-play data and action calls with cards. Output wording varies → assess the decision and evidence, not exact phrasing. Retrieval behaves differently → use the frozen trace. Discussion becomes technical → return to value, authority, reversibility, and accountability.

FINAL EXIT QUESTION: “A proposed business-development agent reads CRM notes, recommends next actions, updates records, and sends commitments to prospects. Would you approve a pilot? State the use-case fit, minimum model/tools/instructions, single- or multi-agent choice, one guardrail, and a human-intervention trigger.” Expected: start with one observable agent, distinguish read/write authority, require approval before commitments, and validate with real users and defined measures.